



AI Agent for Playing Gomoku


Bo Yu

SUNet ID: boyu19

Stanford
Computer Science

Problem Definition

- What is Gomoku?
 - A two-player turn-based zero-sum game
 - Played with Go pieces on a Go board
 - Winner is the player who can first obtain an unbroken chain of five pieces horizontally, vertically, or diagonally
 - White wins in the following example
- Challenges
 - Search space is huge (in a 9×9 board, there are around 80 moves and in a 15×15 board, there are around 225 moves at each state)
 - Good heuristic evaluation function requires deep domain knowledge of the game
- Scope
 - Build AI agents for playing Gomoku based on minimax and Monte Carlo tree search




Model

- Two-player, perfect-info, zero-sum game
 - Players** = {black, white} (black moves first)
 - State s** : (piece positions, whose turn it is)
 - Actions(s)**: all legal positions that Player(s) can place the piece
 - Succ(s, a)**: resulting game state if choosing action a in state s
 - IsEnd(s)**: whether s results in a five-in-a-row or draw
 - Utility(s)**: $+\infty$ if white wins, $-\infty$ if black wins
 - Player(s)** \in **Players**: player who controls state s

Minimax with alpha-beta Pruning

- Limited depth minimax with alpha-beta pruning time complexity
 - $O(b^{2d})$ for branching factor b and depth d
 - For 15×15 board: $b \approx 225$, $b^2 \approx 50625$ ($d = 1$), $b^4 \approx 2562890625$ ($d = 2$)
 - For 9×9 board: $b \approx 81$, $b^2 \approx 6561$ ($d = 1$), $b^4 \approx 43046721$ ($d = 2$)
 - $d = 2$ is prohibitively complex
- Minimax with search space reduction
 - Not all legal actions need to be searched
 - We can reduce the branching factor using certain heuristics (like beam search) and then increase the depth of search tree
- Evaluate function
 - Features: different threat patterns
$$\text{Eval}(s) = 6000 \cdot (N_5^{\text{agent}} - N_5^{\text{opp}}) + 4800 \cdot (N_{\text{open4}}^{\text{agent}} - N_{\text{open4}}^{\text{opp}}) + 500 \cdot (N_{\text{half4}}^{\text{agent}} - N_{\text{half4}}^{\text{opp}}) + 500 \cdot (N_{\text{open3}}^{\text{agent}} - N_{\text{open3}}^{\text{opp}}) + 200 \cdot (N_{\text{half3}}^{\text{agent}} - N_{\text{half3}}^{\text{opp}}) + 50 \cdot (N_{\text{open2}}^{\text{agent}} - N_{\text{open2}}^{\text{opp}}) + 10 \cdot (N_{\text{half2}}^{\text{agent}} - N_{\text{half2}}^{\text{opp}})$$

Monte Carlo Tree Search



- Tree policy: using UCB1 algorithm
- Default policy: roll out simulations to completion
 - Random policy
 - Incorporate domain knowledge

Experiments and Results

- Environment: 9×9 and 15×15 board
- Agents
 - Baseline:
 - Greedy heuristic algorithm: randomly select one of the top K candidates according to the heuristic board evaluation function Eval(s)
 - Minimax with alpha-beta pruning:
 - Depth = 1 without search space reduction
 - Depth = 2 with search space reduction ($b = 5/10$)
 - Monte Carlo tree search:
 - Random roll-out policy with 1000/2000/4000/8000 sim. budget
 - Heuristic roll-out policy with 200/500 sim. budget
- Comparison with baseline:

Agents	Winning ratio	Average time per step (sec)
Minimax-d1	20/20	0.73
Minimax-d2 ($b = 5$)	20/20	3.3
MCTS-random-1000	7/20	15.1
MCTS-random-2000	9/20	45.5
MCTS-random-4000	13/20	105.1
MCTS-random-8000	20/20	140.7
MCTS-heuristic-200	8/20	90.3
MCTS-heuristic-500	13/20	191.3

- Comparison with minimax depth = 1

9×9 board	Winning ratio	Average time per step (sec)
Minimax-d2 ($b = 5$)	6W/10D/4L	3.9
Minimax-d2($b = 10$)	1W/9D	16.2

15×15 board	Winning ratio	Average time per step (sec)
Minimax-d2 ($b = 5$)	20L	10.7
Minimax-d2($b = 10$)	6W/3L	58.3

Analysis

- Minimax algorithm:
 - Depth = 1:
 - Performs the best considering both winning percentage and time
 - Depth = 2:
 - Need to reduce branching factor significantly, which degrades accuracy and performance
- Monte Carlo tree search algorithm:
 - With random roll-out policy:
 - No domain knowledge or handcrafted evaluation functions are needed
 - Need large number of random simulations (self-plays) to get good performance (converge to minimax policy), which is very time-consuming
 - With heuristic roll-out policy:
 - Incorporate some domain knowledge during roll-out stage instead of using random policy
 - Require a smaller number of simulations, but each simulation takes longer time

Future Work

- Minimax algorithm:
 - More sophisticated features for the evaluation function based on expert knowledge of the game
 - Use TD learning to learn the weights of the evaluation function
- Monte Carlo tree search algorithm:
 - Use deep neural networks to guide MCTS to reduce the effective depth and breadth of the search tree, thus improve the efficiency and accuracy of the algorithm

References

- [1] Sylvain Gelly, David Silver, "Monte-Carlo tree search and rapid action value estimation in computer Go," Artificial Intelligence, vol. 175, Issue 11, 2011, pp. 1856-1875.
- [2] C. B. Browne *et al.*, "A Survey of Monte Carlo Tree Search Methods," in IEEE Transactions on Computational Intelligence and AI in Games, vol. 4, no. 1, pp. 1-43, March 2012.